

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv("ionosphere.csv")
df
```

```
Out[2]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.511
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.265
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.402
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.906
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.651
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.015
...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.042
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.013
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.031
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.020
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.151

350 rows × 35 columns



```
In [3]: df.head()
```

```
Out[3]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.51171
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535

5 rows × 35 columns



Data Cleaning and Data Preprocessing

In [4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 35 columns):
#   Column          Non-Null Count  Dtype
---  -
0   1                350 non-null   int64
1   0                350 non-null   int64
2   0.99539          350 non-null   float64
3   -0.05889         350 non-null   float64
4   0.85243          350 non-null   float64
5   0.02306          350 non-null   float64
6   0.83398          350 non-null   float64
7   -0.37708         350 non-null   float64
8   1.1              350 non-null   float64
9   0.03760          350 non-null   float64
10  0.85243.1        350 non-null   float64
11  -0.17755         350 non-null   float64
12  0.59755          350 non-null   float64
13  -0.44945         350 non-null   float64
14  0.60536          350 non-null   float64
15  -0.38223         350 non-null   float64
16  0.84356          350 non-null   float64
17  -0.38542         350 non-null   float64
18  0.58212          350 non-null   float64
19  -0.32192         350 non-null   float64
20  0.56971          350 non-null   float64
21  -0.29674         350 non-null   float64
22  0.36946          350 non-null   float64
23  -0.47357         350 non-null   float64
24  0.56811          350 non-null   float64
25  -0.51171         350 non-null   float64
26  0.41078          350 non-null   float64
27  -0.46168         350 non-null   float64
28  0.21266          350 non-null   float64
29  -0.34090         350 non-null   float64
30  0.42267          350 non-null   float64
31  -0.54487         350 non-null   float64
32  0.18641          350 non-null   float64
33  -0.45300         350 non-null   float64
34  g                350 non-null   object
dtypes: float64(32), int64(2), object(1)
memory usage: 95.8+ KB
```

In [5]: `df.describe()`

Out[5]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708
count	350.000000	350.0	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000
mean	0.891429	0.0	0.640330	0.044667	0.600350	0.116154	0.549284	0.120779
std	0.311546	0.0	0.498059	0.442032	0.520431	0.461443	0.493124	0.520816
min	0.000000	0.0	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000
25%	1.000000	0.0	0.471517	-0.065388	0.412555	-0.024868	0.209105	-0.053485
50%	1.000000	0.0	0.870795	0.016700	0.808620	0.021170	0.728000	0.015085
75%	1.000000	0.0	1.000000	0.194727	1.000000	0.335317	0.970445	0.451572
max	1.000000	0.0	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 34 columns



In [6]: `df.columns`

Out[6]: Index(['1', '0', '0.99539', '-0.05889', '0.85243', '0.02306', '0.83398',
'-0.37708', '1.1', '0.03760', '0.85243.1', '-0.17755', '0.59755',
'-0.44945', '0.60536', '-0.38223', '0.84356', '-0.38542', '0.58212',
'-0.32192', '0.56971', '-0.29674', '0.36946', '-0.47357', '0.56811',
'-0.51171', '0.41078', '-0.46168', '0.21266', '-0.34090', '0.42267',
'-0.54487', '0.18641', '-0.45300', 'g'],
dtype='object')

In [7]: `feature_matrix = df.iloc[:,0:34]`
`target_vector = df.iloc[:, -1]`

In [8]: `fs = StandardScaler().fit_transform(feature_matrix)`
`logr = LogisticRegression()`
`logr.fit(fs, target_vector)`

Out[8]: `LogisticRegression()`

```
In [9]: observation=[[1.0,0.0,1.0,-0.18829,0.93035,  
-0.36156,  
-0.10868,  
-0.93597,  
1.0,  
-0.04549,  
0.50874,  
-0.67743,  
0.34432,  
-0.69707,  
-0.51685,  
-0.97515,  
0.05499,  
-0.62237,  
0.33109,  
-1.0,  
-0.13151,  
-0.453,  
-0.18056,  
-0.35734,  
-0.20332,  
-0.26569,  
-0.20468,  
-0.18401,  
-0.1904,  
-0.11593,  
-0.16626,  
-0.06288,  
-0.13738,  
-0.02447]]  
prediction = logr.predict(observation)  
print(prediction)
```

```
['g']
```

```
In [10]: logr.classes_
```

```
Out[10]: array(['b', 'g'], dtype=object)
```

```
In [11]: logr.predict_proba(observation)
```

```
Out[11]: array([[0.07006552, 0.92993448]])
```

```
In [12]: df['g'].value_counts()
```

```
Out[12]: g    224  
b    126  
Name: g, dtype: int64
```

```
In [13]: x=df.drop('g', axis=1)  
y=df['g']
```

```
In [14]: g1={"g":{"g":1, "b":2}}
df=df.replace(g1)
df
```

```
Out[14]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.511
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.265
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.402
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.906
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.651
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.015
...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.042
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.013
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.031
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.020
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.151

350 rows × 35 columns

```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[16]: RandomForestClassifier()
```

```
In [20]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25], 'n_estimators': [10, 20, 30, 40, 50]}
```

```
In [21]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid= parameters,cv=2,scoring =
grid_search.fit(x_train,y_train)
```

```
Out[21]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 3, 4, 5],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```

```
In [22]: grid_search.best_score_
```

```
Out[22]: 0.9221311475409837
```

```
In [25]: rfc_best = grid_search.best_estimator_
```

```
In [27]: from sklearn.tree import plot_tree  
plt.figure(figsize = (80,40))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names = ['Yes',
```

```

Out[27]: [Text(2896.2857142857147, 1993.2, '0.41078 <= 1.0\ngini = 0.454\nsamples = 15
7\nvalue = [85, 159]\nclass = No'),
  Text(2178.857142857143, 1630.8000000000002, '0.18641 <= 0.292\ngini = 0.309
\nsamples = 117\nvalue = [35, 148]\nclass = No'),
  Text(1381.7142857142858, 1268.4, '-0.47357 <= 0.392\ngini = 0.478\nsamples =
49\nvalue = [28, 43]\nclass = No'),
  Text(850.2857142857143, 906.0, '0.56971 <= 0.031\ngini = 0.5\nsamples = 38\n
value = [27, 27]\nclass = Yes'),
  Text(425.14285714285717, 543.5999999999999, '-0.45300 <= -0.003\ngini = 0.40
4\nsamples = 25\nvalue = [23, 9]\nclass = Yes'),
  Text(212.57142857142858, 181.19999999999982, 'gini = 0.497\nsamples = 11\nva
lue = [6, 7]\nclass = No'),
  Text(637.7142857142858, 181.19999999999982, 'gini = 0.188\nsamples = 14\nval
ue = [17, 2]\nclass = Yes'),
  Text(1275.4285714285716, 543.5999999999999, '-0.47357 <= -0.265\ngini = 0.29
8\nsamples = 13\nvalue = [4, 18]\nclass = No'),
  Text(1062.857142857143, 181.19999999999982, 'gini = 0.0\nsamples = 7\nvalue
= [0, 15]\nclass = No'),
  Text(1488.0, 181.19999999999982, 'gini = 0.49\nsamples = 6\nvalue = [4, 3]\n
class = Yes'),
  Text(1913.1428571428573, 906.0, '0.85243 <= 0.862\ngini = 0.111\nsamples = 1
1\nvalue = [1, 16]\nclass = No'),
  Text(1700.5714285714287, 543.5999999999999, 'gini = 0.278\nsamples = 5\nvalu
e = [1, 5]\nclass = No'),
  Text(2125.714285714286, 543.5999999999999, 'gini = 0.0\nsamples = 6\nvalue =
[0, 11]\nclass = No'),
  Text(2976.0, 1268.4, '-0.05889 <= -0.044\ngini = 0.117\nsamples = 68\nvalue
= [7, 105]\nclass = No'),
  Text(2763.4285714285716, 906.0, '0.60536 <= 0.664\ngini = 0.334\nsamples = 1
9\nvalue = [7, 26]\nclass = No'),
  Text(2550.857142857143, 543.5999999999999, 'gini = 0.469\nsamples = 5\nvalue
= [5, 3]\nclass = Yes'),
  Text(2976.0, 543.5999999999999, '-0.32192 <= -0.188\ngini = 0.147\nsamples =
14\nvalue = [2, 23]\nclass = No'),
  Text(2763.4285714285716, 181.19999999999982, 'gini = 0.0\nsamples = 7\nvalue
= [0, 14]\nclass = No'),
  Text(3188.571428571429, 181.19999999999982, 'gini = 0.298\nsamples = 7\nvalu
e = [2, 9]\nclass = No'),
  Text(3188.571428571429, 906.0, 'gini = 0.0\nsamples = 49\nvalue = [0, 79]\nc
lass = No'),
  Text(3613.714285714286, 1630.8000000000002, '0.83398 <= 0.789\ngini = 0.296
\nsamples = 40\nvalue = [50, 11]\nclass = Yes'),
  Text(3401.1428571428573, 1268.4, 'gini = 0.0\nsamples = 23\nvalue = [34, 0]
\nclass = Yes'),
  Text(3826.2857142857147, 1268.4, '-0.47357 <= -0.633\ngini = 0.483\nsamples
= 17\nvalue = [16, 11]\nclass = Yes'),
  Text(3613.714285714286, 906.0, 'gini = 0.0\nsamples = 6\nvalue = [9, 0]\ncla
ss = Yes'),
  Text(4038.857142857143, 906.0, '-0.17755 <= 0.098\ngini = 0.475\nsamples = 1
1\nvalue = [7, 11]\nclass = No'),
  Text(3826.2857142857147, 543.5999999999999, 'gini = 0.375\nsamples = 6\nvalu
e = [2, 6]\nclass = No'),
  Text(4251.428571428572, 543.5999999999999, 'gini = 0.5\nsamples = 5\nvalue =
[5, 5]\nclass = Yes')]

```


