

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("bmi.csv").dropna()
df
```

```
Out[3]:
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns

```
In [4]: df.dropna(inplace=True)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 500 entries, 0 to 499
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Gender  500 non-null      object
1   Height  500 non-null      int64
2   Weight  500 non-null      int64
3   Index   500 non-null      int64
dtypes: int64(3), object(1)
memory usage: 19.5+ KB
```

```
In [6]: feature_matrix = df[['Height','Weight','Index']]
target_vector = df['Gender']
```


3/6

```
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Male', 'Female', 'Female', 'Female', 'Female',
'Female', 'Female', 'Female'], dtype=object)
```

```
In [19]: logr.classes_
```

Out[19]: array(['Female', 'Male'], dtype=object)

```
In [20]: logr.predict_proba(observation)[0][1]
```

Out[20]: 0.013560340101419254

```
In [27]: df['Gender'].value_counts()
```

Out[27]: Female 255
Male 245
Name: Gender, dtype: int64

```
In [30]: x=df.drop('Gender', axis=1)
y=df['Gender']
```

```
In [31]: g1={"Gender":{"Male":1, "Female":2}}
df=df.replace(g1)
df
```

Out[31]:

	Gender	Height	Weight	Index
0	1	174	96	4
1	1	189	87	2
2	2	185	110	4
3	2	195	104	3
4	1	149	61	3
...
495	2	150	153	5
496	2	184	121	4
497	2	141	136	5
498	1	150	95	5

	Gender	Height	Weight	Index
499	1	173	131	5

500 rows × 4 columns

```
In [32]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [33]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[33]: RandomForestClassifier()

```
In [34]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25], 'n_estimators'
```

```
In [35]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid= parameters,cv=2,scoring = "accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[35]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [36]: grid_search.best_score_
```

Out[36]: 0.5057142857142858

```
In [37]: rfc_best = grid_search.best_estimator_
```

```
In [38]: from sklearn.tree import plot_tree
plt.figure(figsize = (80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names = ['Yes','No'],fi
```

```
Out[38]: [Text(2678.3999999999996, 1812.0, 'Weight <= 137.5\ngini = 0.499\nsamples = 216\nvalue =
[168, 182]\nnclass = No'),
Text(1785.6, 1087.2, 'Index <= 1.5\ngini = 0.498\nsamples = 163\nvalue = [138, 123]\nnclass = Yes'),
Text(892.8, 362.399999999999986, 'gini = 0.465\nsamples = 13\nvalue = [7, 12]\nnclass = No'),
Text(2678.3999999999996, 362.399999999999986, 'gini = 0.497\nsamples = 150\nvalue = [131, 111]\nnclass = Yes'),
Text(3571.2, 1087.2, 'gini = 0.447\nsamples = 53\nvalue = [30, 59]\nnclass = No')]
```

